

Understanding TF-IDF (Term Frequency-Inverse Document Frequency)

Last Updated : 07 Feb, 2025

TF-IDF (Term Frequency-Inverse Document Frequency) is a statistical measure used in natural language processing and information retrieval **to evaluate the importance of a word in a document relative to a collection of documents (corpus).**

Unlike simple word frequency, TF-IDF balances common and rare words to highlight the most meaningful terms.

How TF-IDF Works?

TF-IDF combines two components: Term Frequency (TF) and Inverse Document Frequency (IDF).

Term Frequency (TF): Measures how often a word appears in a document. A higher frequency suggests greater importance. If a term appears frequently in a document, it is likely relevant to the document's content. **Formula:**

$$TF(t, d) = \frac{\text{Number of times term } t \text{ appears in document } d}{\text{Total number of terms in document } d}$$

Term Frequency (TF)

Limitations of TF Alone:

- TF does not account for the global importance of a term across the entire corpus.
- Common words like **"the"** or **"and"** may have high TF scores but are not meaningful in distinguishing documents.

appears in fewer documents, it is more likely to be meaningful and specific.

Formula:

$$\text{IDF}(t, D) = \log \frac{\text{Total number of documents in corpus } D}{\text{Number of documents containing term } t}$$

Inverse Document Frequency (IDF)

- The logarithm is used to dampen the effect of very large or very small values, ensuring the IDF score scales appropriately.
- It also helps balance the impact of terms that appear in extremely few or extremely many documents.

Limitations of IDF Alone:

- IDF does not consider how often a term appears within a specific document.
- A term might be rare across the corpus (high IDF) but irrelevant in a specific document (low TF).

Converting Text into vectors with TF-IDF : Example

To better grasp how TF-IDF works, let's walk through a detailed example.

Imagine we have a **corpus** (a collection of documents) with three documents:

1. **Document 1:** "The cat sat on the mat."
2. **Document 2:** "The dog played in the park."
3. **Document 3:** "Cats and dogs are great pets."

Our goal is to calculate the TF-IDF score for specific terms in these documents.

Let's focus on the word "**cat**" and see how TF-IDF evaluates its importance.

Step 1: Calculate Term Frequency (TF)

For Document 1:

- The word "**cat**" appears **1 time**.
- The total number of terms in Document 1 is **6** ("the", "cat", "sat", "on", "the", "mat").

For Document 2:

- The word "**cat**" does **not appear**.
- So, $TF(cat, Document\ 2) = 0$.

For Document 3:

- The word "**cat**" appears **1 time** (as "cats").
 - The total number of terms in Document 3 is **6** ("cats", "and", "dogs", "are", "great", "pets").
 - So, $TF(cat, Document\ 3) = 1/6$
-
- *In Document 1 and Document 3, the word "**cat**" has the same TF score. This means it appears with the same relative frequency in both documents.*
 - *In Document 2, the TF score is 0 because the word "**cat**" does not appear.*

Step 2: Calculate Inverse Document Frequency (IDF)

- **Total number of documents in the corpus (D): 3**
- **Number of documents containing the term "cat": 2** (Document 1 and Document 3).

So, $IDF(cat, D) = \log \frac{3}{2} \approx 0.176$

*The IDF score for "**cat**" is relatively low. This indicates that the word "**cat**" is not very rare in the corpus—it appears in 2 out of 3 documents. If a term appeared in only 1 document, its IDF score would be higher, indicating greater uniqueness.*

Step 3: Calculate TF-IDF

The TF-IDF score for "**cat**" is 0.029 in Document 1 and Document 3, and 0 in Document 2 that **reflects both the frequency of the term in the document (TF) and its rarity across the corpus (IDF)**

The TF-IDF score is the product of TF and IDF:

$$\text{TF-IDF}(t, d, D) = \text{TF}(t, d) \times \text{IDF}(t, D)$$

For Document 1:

$$\text{TF-IDF}(\text{cat}, \text{Document 1}, D) = 0.167 \times 0.176 \approx 0.029$$

For Document 2:

$$\text{TF-IDF}(\text{cat}, \text{Document 2}, D) = 0 \times 0.176 = 0$$

For Document 3:

$$\text{TF-IDF}(\text{cat}, \text{Document 3}, D) = 0.167 \times 0.176 \approx 0.029$$

TF-IDF

A higher TF-IDF score means the term is more important in that specific document.

Why is TF-IDF Useful in This Example?

1. Identifying Important Terms: TF-IDF helps us understand that "cat" is somewhat important in Document 1 and Document 3 but irrelevant in Document 2.

If we were building a search engine, this score would help rank Document 1 and Document 3 higher for a query like "cat".

2. Filtering Common Words: Words like "the" or "and" would have high TF scores but very low IDF scores because they appear in almost all documents. Their TF-IDF scores would be close to 0, indicating they are not meaningful.

3. Highlighting Unique Terms: If a term like "mat" appeared only in Document 1, it would have a higher IDF score, making its TF-IDF score more significant in that document.

Implementing TF-IDF in Sklearn with Python

In python tf-idf values can be computed using `TfidfVectorizer()` method in `sklearn` module.

Parameters:

- **input:** It refers to parameter document passed, it can be a filename, file or content itself.

Attributes:

- **vocabulary_:** It returns a dictionary of terms as keys and values as feature indices.
- **idf_:** It returns the inverse document frequency vector of the document passed as a parameter.

Returns:

- **fit_transform():** It returns an array of terms along with tf-idf values.
- **get_feature_names():** It returns a list of feature names.

Step-by-step Approach:

- Import modules.

```
# import required module
from sklearn.feature_extraction.text import TfidfVectorizer
```

- Collect strings from documents and create a corpus having a collection of strings from the documents *d0*, *d1*, and *d2*.

```
# assign documents
d0 = 'Geeks for geeks'
d1 = 'Geeks'
d2 = 'r2j'

# merge documents into a single corpus
string = [d0, d1, d2]
```

- Get tf-idf values from *fit_transform()* method.

```
# create object
tfidf = TfidfVectorizer()
```

- Display idf values of the words present in the corpus.

```
# get idf values
print('\nidf values:')
for ele1, ele2 in zip(tfidf.get_feature_names(), tfidf.idf_):
    print(ele1, ': ', ele2)
```

Output:

```
idf values:
for : 1.6931471805599454
geeks : 1.2876820724517808
r2j : 1.6931471805599454
```

- Display tf-idf values along with indexing.

```
# get indexing
print('\nWord indexes:')
print(tfidf.vocabulary_)

# display tf-idf values
print('\ntf-idf value:')
print(result)

# in matrix form
print('\ntf-idf values in matrix form:')
print(result.toarray())
```

Output:

```
Word indexes:
{'geeks': 1, 'for': 0, 'r2j': 2}
```

```
tf-idf value:
(0, 0)      0.5493512310263033
(0, 1)      0.8355915419449176
(1, 1)      1.0
(2, 2)      1.0
```

```
tf-idf values in matrix form:
[[0.54935123 0.83559154 0.
  [0.          1.          0.          ]
  [0.          0.          1.          ]]
```

```
Word indexes:
{'geeks': 1, 'for': 0, 'r2j': 2}
```

```
tf-idf value:
(0, 0)    0.5493512310263033
(0, 1)    0.8355915419449176
(1, 1)    1.0
(2, 2)    1.0
```

tf-idf value of word having index 1 i.e. *geeks* in document index 0 i.e. *d0*

Document Index

Word Index

From the above image the below table can be generated:

Document	Word	Document Index	Word Index	tf-idf value
d0	for	0	0	0.549
d0	geeks	0	1	0.8355
d1	geeks	1	1	1.000
d2	r2j	2	2	1.000

Below are some examples which depict how to compute tf-idf values of words from a corpus:

Example 1: Below is the complete program based on the above approach:

```
# import required module
from sklearn.feature_extraction.text import TfidfVectorizer

# assign documents
d0 = 'Geeks for geeks'
d1 = 'Geeks'
d2 = 'r2j'

# merge documents into a single corpus
string = [d0, d1, d2]

# create object
tfidf = TfidfVectorizer()
```

```

# get idf values
print('\nidf values:')
for ele1, ele2 in zip(tfidf.get_feature_names(), tfidf.idf_):
    print(ele1, ': ', ele2)

# get indexing
print('\nWord indexes:')
print(tfidf.vocabulary_)

# display tf-idf values
print('\ntf-idf value:')
print(result)

# in matrix form
print('\ntf-idf values in matrix form:')
print(result.toarray())

```

Output:

```

idf values:
for : 1.6931471805599454
geeks : 1.2876820724517808
r2j : 1.6931471805599454

Word indexes:
{'geeks': 1, 'for': 0, 'r2j': 2}

tf-idf value:
(0, 0)      0.5493512310263033
(0, 1)      0.8355915419449176
(1, 1)      1.0
(2, 2)      1.0

tf-idf values in matrix form:
[[0.54935123 0.83559154 0.
  [0.          1.          0.          ]
  [0.          0.          1.          ]]

```

Example 2: Here, tf-idf values are computed from a corpus having unique values.

```

# import required module
from sklearn.feature_extraction.text import TfidfVectorizer

# assign documents
d0 = 'geek1'
d1 = 'geek2'
d2 = 'geek3'
d3 = 'geek4'

# merge documents into a single corpus
string = [d0, d1, d2, d3]

```




```

# get tf-df values
result = tfidf.fit_transform(string)

# get indexing
print('\nWord indexes:')
print(tfidf.vocabulary_)

# display tf-idf values
print('\ntf-idf values:')
print(result)

```

Output:

```

Word indexes:
{'geek1': 0, 'geek2': 1, 'geek3': 2, 'geek4': 3}

tf-idf values:
(0, 0)      1.0
(1, 1)      1.0
(2, 2)      1.0
(3, 3)      1.0

```

Example 3: In this program, tf-idf values are computed from a corpus having similar documents.

```

# import required module
from sklearn.feature_extraction.text import TfidfVectorizer

# assign documents
d0 = 'Geeks for geeks!'
d1 = 'Geeks for geeks!'

# merge documents into a single corpus
string = [d0, d1]

# create object
tfidf = TfidfVectorizer()

# get tf-df values
result = tfidf.fit_transform(string)

# get indexing
print('\nWord indexes:')
print(tfidf.vocabulary_)

# display tf-idf values
print('\ntf-idf values:')

```

Output:

```
Word indexes:
{'geeks': 1, 'for': 0}

tf-idf values:
(0, 0)      0.4472135954999579
(0, 1)      0.8944271909999159
(1, 0)      0.4472135954999579
(1, 1)      0.8944271909999159
```

Example 4: Below is the program in which we try to calculate tf-idf value of a single word *geeks* is repeated multiple times in multiple documents.

```
# import required module
from sklearn.feature_extraction.text import TfidfVectorizer

# assign corpus
string = ['Geeks geeks']*5

# create object
tfidf = TfidfVectorizer()

# get tf-df values
result = tfidf.fit_transform(string)

# get indexing
print('\nWord indexes:')
print(tfidf.vocabulary_)

# display tf-idf values
print('\ntf-idf values:')
print(result)
```

Output:

```
Word indexes:
{'geeks': 0}

tf-idf values:
(0, 0)      1.0
(1, 0)      1.0
(2, 0)      1.0
(3, 0)      1.0
```

[Comment](#)[More info](#)[Advertise with us](#)

Next Article

Counting Word Frequency and
Making a Dictionary from it

Similar Reads

Bag of word and Frequency count in text using sklearn

Text data is ubiquitous in today's digital world, from emails and social media posts to research articles and customer reviews. To analyze and derive insights from this textual information, it's essential to convert text into...

3 min read

Understanding min_df and max_df in scikit CountVectorizer

In natural language processing (NLP), text preprocessing is a critical step that significantly impacts the performance of machine learning models. One common preprocessing step is converting text data into...

6 min read

Sort elements by frequency | Set 5 (using Java Map)

Given an integer array, sort the array according to the frequency of elements in decreasing order, if the frequency of two elements are same then sort in increasing order Examples: Input: arr[] = {2, 3, 2, 4, 5, 12, 2, ...

3 min read

Counting Word Frequency and Making a Dictionary from it

We need to count how often each word appears in a given text and store these counts in a dictionary. For instance, if the input string is "Python with Python gfg with Python", we want the output to be {'Python': 3,...

3 min read

Classification of text documents using sparse features in Python Scikit Learn

Classification is a type of machine learning algorithm in which the model is trained, so as to categorize or label the given input based on the provided features for example classifying the input image as an image of a dog o...

5 min read

How to Fix "Document-Term Matrix in R - Bigram Tokenizer Not Working"?

Creating a Document-Term Matrix (DTM) is crucial in text analysis, allowing you to explore the frequency of terms within a set of documents. Sometimes, using a Bigram tokenizer with your DTM in R may lead to

An Easy Approach to TF-IDF Using R

TF-IDF (Term Frequency-Inverse Document Frequency) is a fundamental technique in natural language processing and information retrieval for assessing the importance of a term within a document relative to a...

5 min read

Document Similarity Using LSA in R

Latent Semantic Analysis (LSA) is a powerful technique in natural language processing (NLP) that helps uncover the hidden relationships between words in a set of documents. Unlike basic word-matching...

5 min read

Sort Dataframe according to row frequency in Pandas

In this article, we will discuss how to use `count()` and `sort_values()` in pandas. So the count in pandas counts the frequency of elements in the dataframe column and then sort sorts the dataframe according to element...

2 min read

Maximum repeated frequency of characters in a given string

Given a string S, the task is to find the count of maximum repeated frequency of characters in the given string S. Examples: Input: S = "geeksgeeks" Output: Frequency 2 is repeated 3 times Explanation: Frequency of...

6 min read



Corporate & Communications Address:

A-143, 7th Floor, Sovereign Corporate
Tower, Sector- 136, Noida, Uttar Pradesh
(201305)

Registered Address:

K 061, Tower K, Gulshan Vivante
Apartment, Sector 137, Noida, Gautam
Buddh Nagar, Uttar Pradesh, 201305



Company

About Us
Legal
Privacy Policy
In Media
Contact Us
Advertise with us
GFG Corporate
Solution
Placement Training
Program

Languages

Python
Java
C++
PHP
GoLang
SQL
R Language
Android Tutorial
Tutorials Archive

DSA

Data Structures
Algorithms
DSA for Beginners
Basic DSA
Problems
DSA Roadmap
Top 100 DSA
Interview Problems
DSA Roadmap by
Sandeep Jain
All Cheat Sheets

Data Science & ML

Data Science With
Python
Data Science For
Beginner
Machine Learning
ML Maths
Data Visualisation
Pandas
NumPy
NLP
Deep Learning

Web Technologies

HTML
CSS
JavaScript
TypeScript
ReactJS
NextJS
Bootstrap
Web Design

Python Tutorial

Python
Programming
Examples
Python Projects
Python Tkinter
Python Web
Scraping
OpenCV Tutorial
Python Interview
Question
Django

Computer Science

Operating Systems
Computer Network
Database
Management
System
Software
Engineering
Digital Logic Design
Engineering Maths
Software
Development
Software Testing

DevOps

Git
Linux
AWS
Docker
Kubernetes
Azure
GCP
DevOps Roadmap

System Design

High Level Design
Low Level Design
UML Diagrams
Interview Guide
Design Patterns
OOAD
System Design
Bootcamp
Interview
Questions

Interview Preparation

Competitive
Programming
Top DS or Algo for
CP
Company-Wise
Recruitment
Process
Company-Wise
Preparation
Aptitude
Preparation
Puzzles

School Subjects

Mathematics
Physics
Chemistry
Biology
Social Science
English Grammar
Commerce
World GK

GeeksforGeeks Videos

DSA
Python
Java
C++
Web Development
Data Science
CS Subjects

@GeeksforGeeks, Sanchhaya Education Private Limited, All rights reserved